

```
!pip install -q keras
import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras import models, datasets, layers
import matplotlib.pyplot as plt
import matplotlib.image as mp
```

```
(train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist11490434/11490434> [=====] - 0s 0us/step



```
print('x_train:', train_images.shape)
print('y_train:', train_labels.shape)
print('x_test:', test_images.shape)
print('y_test:', test_labels.shape)
```

```
x_train: (60000, 28, 28)
y_train: (60000,)
x_test: (10000, 28, 28)
y_test: (10000,)
```

```
pd.DataFrame(train_images[0])
```

	0	1	2	3	4	5	6	7	8	9	...	18	19	20	21	22
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	...	175	26	166	255	247
6	0	0	0	0	0	0	0	0	30	36	...	225	172	253	242	195
7	0	0	0	0	0	0	0	49	238	253	...	93	82	82	56	39
8	0	0	0	0	0	0	0	18	219	253	...	0	0	0	0	0
9	0	0	0	0	0	0	0	0	80	156	...	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	14	...	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
12	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
13	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
14	0	0	0	0	0	0	0	0	0	0	...	25	0	0	0	0

```
train_images=train_images/255
```

```
test_images= test_images/255
```

```
model=models.Sequential()
model.add(layers.Flatten(input_shape=(28, 28, 1)))
model.add(layers.Dense(32,activation="relu"))
model.add(layers.Dense(16,activation="relu"))
model.add(layers.Dense(10,activation="softmax"))
model.compile(optimizer='adam',loss="sparse_categorical_crossentropy",metrics=['accuracy'])
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 784)	0
dense_9 (Dense)	(None, 32)	25120
dense_10 (Dense)	(None, 16)	528
dense_11 (Dense)	(None, 10)	170


```
=====
Total params: 25,818
Trainable params: 25,818
Non-trainable params: 0
=====
```

```
h1 = model.fit(train_images, train_labels, epochs=10, validation_data=(test_images, test_labels))
```

```
Epoch 1/10
1875/1875 [=====] - 7s 3ms/step - loss: 0.3697 - accuracy: 0.9000
Epoch 2/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.1702 - accuracy: 0.9500
Epoch 3/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.1365 - accuracy: 0.9600
Epoch 4/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.1182 - accuracy: 0.9650
Epoch 5/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.1043 - accuracy: 0.9700
Epoch 6/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.0944 - accuracy: 0.9750
Epoch 7/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0851 - accuracy: 0.9780
Epoch 8/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.0776 - accuracy: 0.9800
Epoch 9/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.0722 - accuracy: 0.9820
Epoch 10/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.0673 - accuracy: 0.9850
```

```
import matplotlib.pyplot as plt
#Plotting the training and validation loss and accuracy
f,ax=plt.subplots(2,1)

#Loss
ax[0].plot(h1.history['loss'], label='train loss')
ax[0].plot(h1.history['val_loss'], label='val loss')
ax[0].legend(loc='upper right')
#Accuracy
ax[1].plot(h1.history['accuracy'],color='b',label='Training Accuracy')
ax[1].plot(h1.history['val_accuracy'],color='r',label='Validation Accuracy')
ax[1].legend(loc='lower right');
```



```
model.compile(optimizer='sgd', loss="sparse_categorical_crossentropy", metrics=['accuracy'])
h2 = model.fit(train_images, train_labels, epochs=10, validation_data=(test_images, test_labels))
```

```

Epoch 1/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.8424 - accuracy: 0.3601
Epoch 2/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.3294 - accuracy: 0.6000
Epoch 3/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.2712 - accuracy: 0.6000
Epoch 4/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.2392 - accuracy: 0.6000
Epoch 5/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.2180 - accuracy: 0.6000
Epoch 6/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.2018 - accuracy: 0.6000
Epoch 7/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.1885 - accuracy: 0.6000
Epoch 8/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.1778 - accuracy: 0.6000
Epoch 9/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.1673 - accuracy: 0.6000
Epoch 10/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.1593 - accuracy: 0.6000

```

```


#Plotting the training and validation loss and accuracy
f,ax=plt.subplots(2,1)

```

```

#Loss
ax[0].plot(h2.history['loss'],color='b',label='Training Loss')
ax[0].plot(h2.history['val_loss'],color='r',label='Validation Loss')
ax[0].legend(loc='center right');
#Accuracy
ax[1].plot(h2.history['accuracy'],color='b',label='Training Accuracy')
ax[1].plot(h2.history['val_accuracy'],color='r',label='Validation Accuracy')
ax[1].legend(loc='center right');

```



```

model.compile(optimizer='rmsprop',loss="sparse_categorical_crossentropy",metrics=['accuracy'])
h3 = model.fit(train_images, train_labels, epochs=10,validation_data=(test_images,test_labels))

```

```

Epoch 1/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.3601 - accuracy: 0.3601
Epoch 2/10

```

```

1875/1875 [=====] - 5s 3ms/step - loss: 0.1956 - accuracy: 0.60
Epoch 3/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.1607 - accuracy: 0.65
Epoch 4/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.1384 - accuracy: 0.70
Epoch 5/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.1234 - accuracy: 0.75
Epoch 6/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.1130 - accuracy: 0.78
Epoch 7/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1050 - accuracy: 0.80
Epoch 8/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0986 - accuracy: 0.82
Epoch 9/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0930 - accuracy: 0.84
Epoch 10/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0909 - accuracy: 0.85

```

```

#Plotting the training and validation loss and accuracy
f,ax=plt.subplots(2,1)

```

```
#Loss
```

```

ax[0].plot(h3.history['loss'],color='b',label='Training Loss')
ax[0].plot(h3.history['val_loss'],color='r',label='Validation Loss')
ax[0].legend(loc='center right');

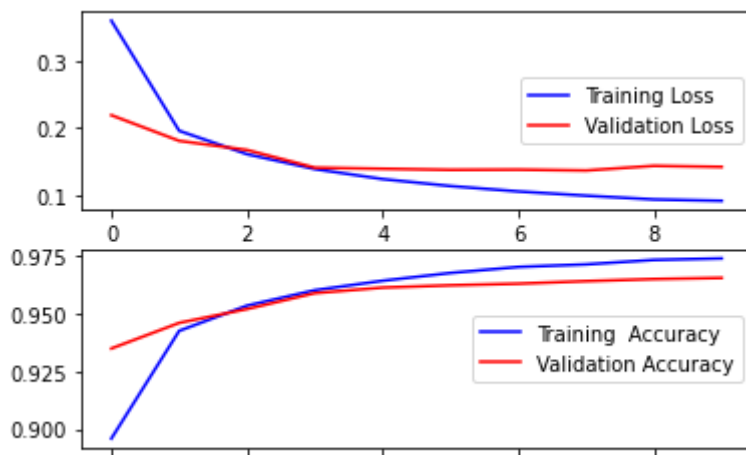
```

```
#Accuracy
```

```

ax[1].plot(h3.history['accuracy'],color='b',label='Training Accuracy')
ax[1].plot(h3.history['val_accuracy'],color='r',label='Validation Accuracy')
ax[1].legend(loc='center right');

```



```

model.compile(optimizer='nadam',loss="sparse_categorical_crossentropy",metrics=['accuracy'])
h4 = model.fit(train_images, train_labels, epochs=10,validation_data=(test_images,test_labels))

```

```

Epoch 1/10
1875/1875 [=====] - 7s 3ms/step - loss: 0.3952 - accuracy: 0.89
Epoch 2/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1808 - accuracy: 0.94
Epoch 3/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1420 - accuracy: 0.95

```

```

Epoch 4/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1202 - accuracy: 0.95
Epoch 5/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1047 - accuracy: 0.96
Epoch 6/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0937 - accuracy: 0.97
Epoch 7/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0862 - accuracy: 0.97
Epoch 8/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0783 - accuracy: 0.98
Epoch 9/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0740 - accuracy: 0.98
Epoch 10/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0675 - accuracy: 0.99

```

```

#Plotting the training and validation loss and accuracy
f,ax=plt.subplots(2,1)

```

```
#Loss
```

```

ax[0].plot(h4.history['loss'],color='b',label='Training Loss')
ax[0].plot(h4.history['val_loss'],color='r',label='Validation Loss');
ax[0].legend(loc='center right');


```

```
#Accuracy
```

```

ax[1].plot(h4.history['accuracy'],color='b',label='Training Accuracy')
ax[1].plot(h4.history['val_accuracy'],color='r',label='Validation Accuracy')
ax[1].legend(loc='center right');

```



```

score = model.evaluate(test_images,test_labels)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

313/313 [=====] - 1s 2ms/step - loss: 0.1198 - accuracy: 0.99
Test score: 0.1198384165763855
Test accuracy: 0.9642999768257141

```

```

model_name = 'digits_recognition.h5'
model.save(model_name, save_format='h5')

```

```
loaded_model = tf.keras.models.load_model(model_name)
predictions_one_hot = loaded_model.predict([test_images])
print('predictions_one_hot:', predictions_one_hot.shape)

313/313 [=====] - 1s 1ms/step
predictions_one_hot: (10000, 10)
```

```
# Let's extract predictions with highest probabilities and detect what digits have been act
predictions = np.argmax(predictions_one_hot, axis=1)
pd.DataFrame(predictions)
plt.imshow(test_images[0].reshape((28,28)), cmap=plt.cm.binary)
plt.show()
```

